

# Kubernetes Networking



## Table of Contents

- ▶ Cluster Networking
- ▶ Services
- ▶ Service Types
- ▶ Labels and loose coupling

# 1 Cluster Networking

## Cluster Networking



## Cluster Networking


There are 4 distinct networking problems to address:

1. container-to-container communications:
2. Pod-to-Pod communications:
3. Pod-to-Service communications: this is covered by services.
4. External-to-Service communications: this is covered by services.




2

## Services




## Services

Pods are not reliable



## Services

A **Service** offers a single **DNS entry** for a containerized application managed by the Kubernetes cluster

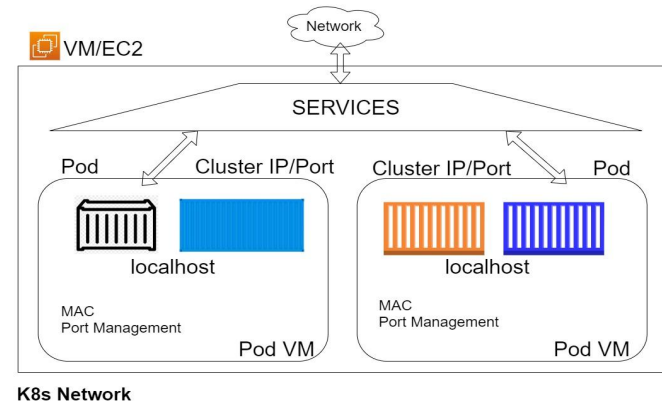



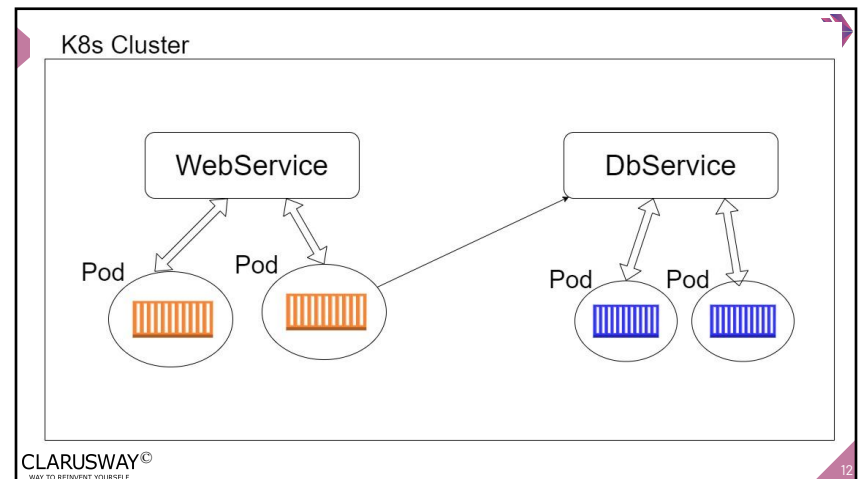
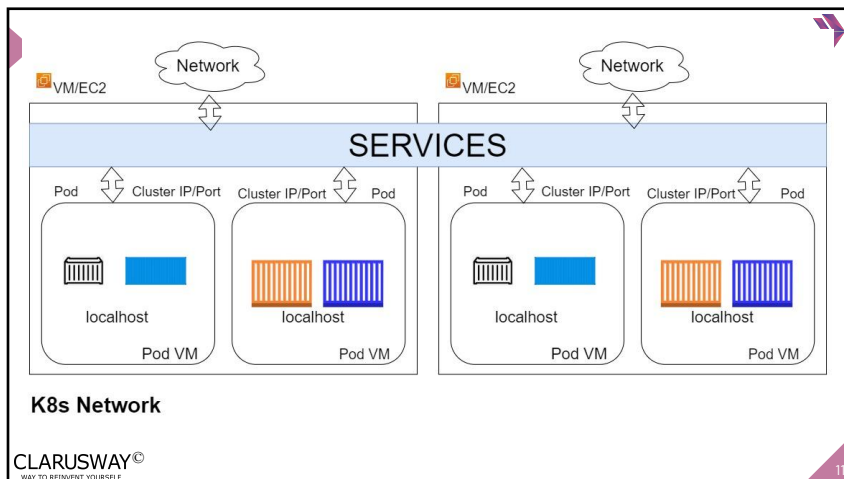
## Services

The **Service** is associated with the Pods, and provides them with a stable IP, DNS and port. It also **loadbalances** requests across the Pods.

**Service** logically groups Pods and defines a policy to access them.

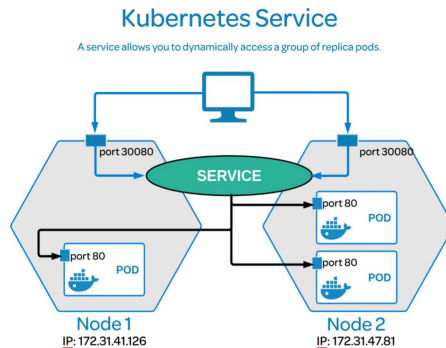
This grouping is achieved via **Labels** and **Selectors**.





## Services

Kubernetes **Services** enable communication between various components **within** and **outside** of the application. Kubernetes Services helps us connect applications together with other applications or users.



## kube-proxy

- Each cluster node runs a daemon called **kube-proxy**
- **kube-proxy** is responsible for **implementing the Service configuration** on behalf of an administrator or developer
- For each new Service, on each node, **kube-proxy** configures **iptables** rules to capture the traffic for its **ClusterIP** and forwards it to one of the Service's endpoints.
- When the Service is removed, **kube-proxy** removes the corresponding **iptables** rules on all nodes as well.

## Service Discovery


- Kubernetes has an add-on for **DNS**, which creates a DNS record for each Service and its format is

`web-svc.my-namespace.svc.cluster.local.`


Service Name    Namespace    Service    domain

- Services within the same Namespace find other Services just by their names.
- If we add a Service **redis-master** in **my-ns** Namespace, all Pods in the same **my-ns** Namespace lookup the Service just by its name, **redis-master**.

**my-ns**



**test-ns**



To connect to the "Game pod" and "db pod":

From "web pod" -> "Game pod" -> hostname: game-service.my-ns:port  
game-service:port

From "web pod" -> "db pod" -> hostname: db-service.test-ns.svc.cluster.local:port




3

## Service Types

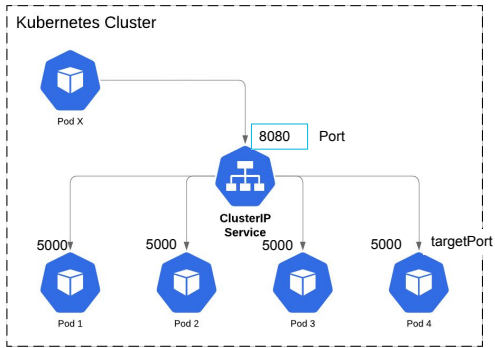
## Service Types

There are 4 major service types:

- ClusterIP (default)
- NodePort
- LoadBalancer
- ExternalName



## Service Types




**ClusterIP:**  
Expose traffic internally

### Example Usecase:


Good for service of  
database & back-end  
apps.

```
apiVersion: v1
kind: Service
metadata:
  name: back-end-svc
  labels:
    app: back-end
spec:
  type: ClusterIP (default)
  selector:
    app: back-end
  ports:
    - port: 8080
      protocol: TCP
      targetPort: 5000
```



Worker Node-1

## Service Types




**NodePort:**  
Exposes traffic to the outside.

### Example Usecase:

when we want to make our Services which has our app or website accessible from the external world.

```
apiVersion: v1
kind: Service
metadata:
  name: web-flask-svc
  labels:
    app: web-flask
spec:
  type: NodePort
  selector:
    app: web-flask
  ports:
    - nodePort: 30480
      port: 3000
      protocol: TCP
      targetPort: 5000
```





Exposes traffic outside with load balancing feature.

when we want to load balancing our Services which has our app or website accessible from the external world.




## Service Types

### LoadBalancer:

- The **LoadBalancer** *ServiceType* will only work if the underlying infrastructure supports the automatic creation of Load Balancers and have the respective support in Kubernetes, as is the case with the Google Cloud Platform, Azure or AWS.
- If no such feature is configured, the **LoadBalancer IP** address field is **not populated**, it remains in **Pending** state, but the **Service will still work as a typical NodePort type Service**.

## Service Types



## Service Types

### ExternalName:

Maps the Service to the contents of the ExternalName field (e.g. example.com), by returning a CNAME record with its value.

### Example Usecases:

to make externally configured services like;

**remote.server.url.com**

available to applications inside the cluster.

## Service Types

```
apiVersion: v1
kind: Service
metadata:
  labels: io.kompose.service: mysql-server
  name: mysql-server
spec:
  type: ExternalName
  externalName:
    serdar.cbanmzptkrzf.us-east-1.rds.amazonaws.com
```

4

## Labels and loose coupling



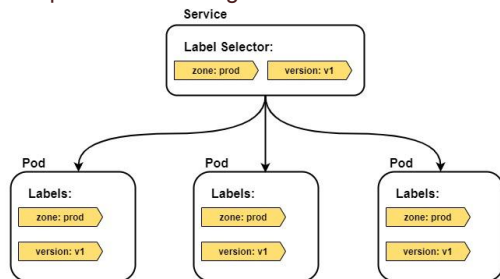
## Labels and loose coupling



- Labels and Selectors use a **key/value** pair format.
- Pods and Services are loosely coupled via labels and label selectors.
- For a Service to match a set of Pods, and therefore provide stable networking and load-balance, it only needs to match some of the Pods labels.
- However, for a Pod to match a Service, the Pod must match all of the values in the Service's label selector.


## Labels and loose coupling

The figure below shows an example where 3 Pods are labeled as **zone=prod** and **version=v1**, and the Service has a label selector that matches. This Service provides stable networking to all three Pods. It also provides simple load-balancing.



## Labels and loose coupling

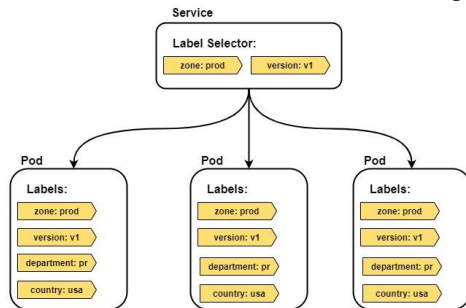
The figure below shows an example where the Service does not match any of the Pods. This is because the Service is selecting on two labels, but the Pods only have one of them. The logic behind this is a Boolean AND operation.





## Labels and loose coupling

This figure shows an example that does work. It doesn't matter that the Pods have additional labels that the Service is not selecting on.



# THANKS!

## Any questions?

You can find me at:

- ▶ Email: [serdar@clarusway.com](mailto:serdar@clarusway.com)
- ▶ Slack: @serdar - instructor



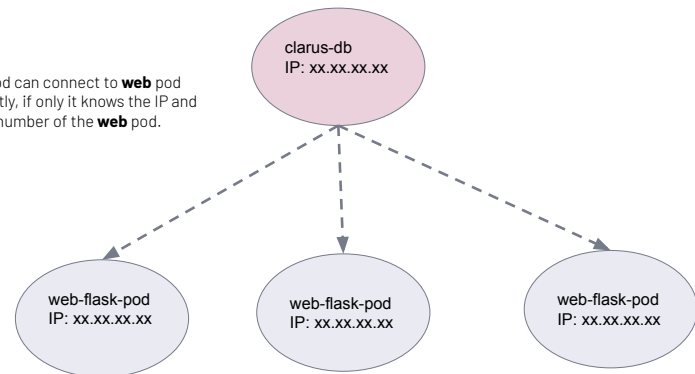
CLARUSWAY®  
Students, write your response!

5

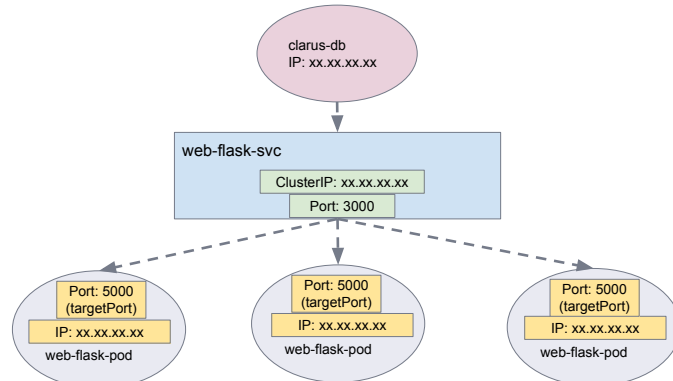
## Kubernetes hands-on-03

## Pod to Pod Connection

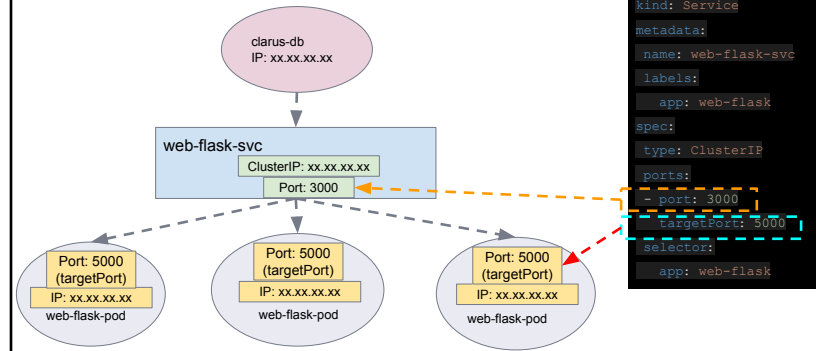
**db** pod can connect to **web** pod directly, if only it knows the IP and port number of the **web** pod.

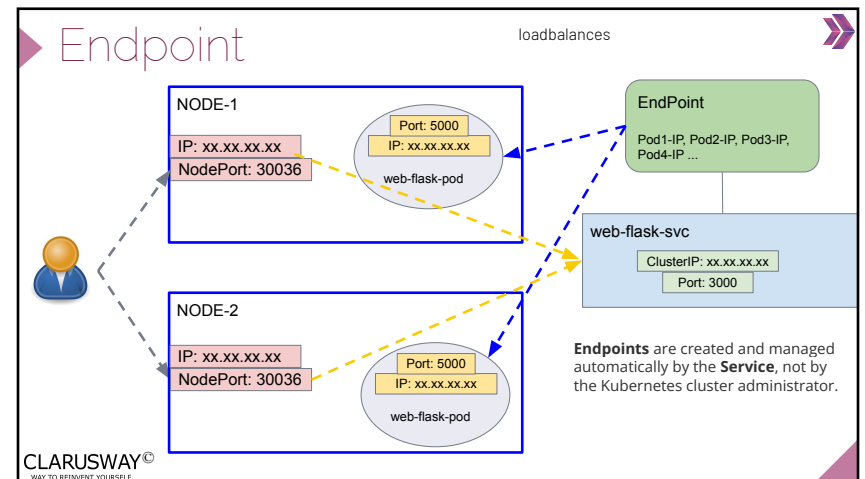
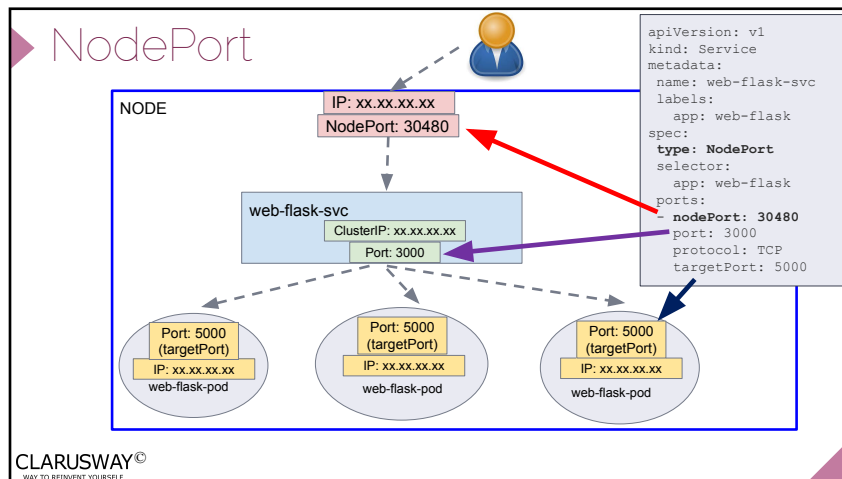


## ClusterIP



## ClusterIP






## Service Discovery

- Pods from other Namespaces, such as `test-ns`, lookup the same Service by adding the respective Namespace as a suffix, such as `redis-master.my-ns` or providing the **FQDN** of the service as


`redis-master.my-ns.svc.cluster.local.`



## Service Types


**NodePort:** Exposes the Service on each Node's IP at a static port (the NodePort). A ClusterIP Service, to which the NodePort Service **routes**, is automatically created. Port can either be **statically** defined, or **dynamically** taken from a range between 30000-32767.

With the **NodePort** *ServiceType*, in addition to a ClusterIP, a high-port is mapped to the respective Service, from all the worker nodes.



## Service Types


**LoadBalancer:** Exposes the Service externally using a cloud provider's load balancer. The external load balancer routes to the automatically created NodePort and ClusterIP Services.



## Service Types

### NodePort:


The end-user connects to any worker node on the specified high-port, which proxies the request internally to the ClusterIP of the Service, then the request is forwarded to the applications running inside the cluster. Let's not forget that the Service is load balancing such requests, and only forwards the request to one of the Pods running the desired application. To manage access to multiple application Services from the external world, administrators can configure a reverse proxy - an ingress, and define rules that target specific Services within the cluster.



## Service Types

### LoadBalancer:

- NodePort and ClusterIP are automatically created, and the external load balancer will route to them
- The Service is exposed at a static port on each worker node
- The Service is exposed externally using the underlying cloud provider's load balancer feature.



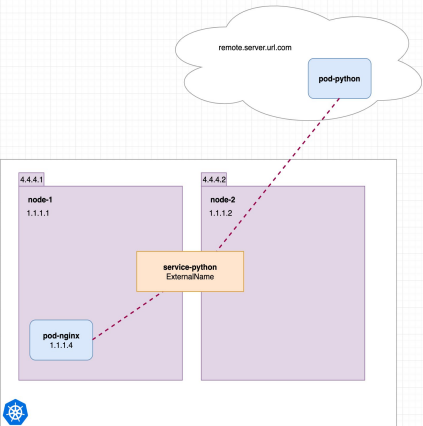
## Service Types

**ExternalName** is a special *ServiceType*, that has no Selectors and does not define any endpoints.

When accessed within the cluster, it returns a **CNAME** record of an externally configured Service.

The primary use case of this *ServiceType* is to make externally configured Services like **my-database.example.com** available to applications inside the cluster.

If the externally defined Service resides within the same Namespace, using just the name **my-database** would make it available to other applications and Services within that same Namespace.



```

kind: Service
apiVersion: v1
metadata:
  name: service-python
spec:
  type: ExternalName
  externalName: remote.server.url.com
  ports:
    - port: 3000
      protocol: TCP
      targetPort: 443

```

Now **pod-nginx** can simply connect to `http://service-python:3000`, just like with a ClusterIP service.

CLARUSWAY  
WAY TO REINVENT YOURSELF

47

## Service Types

```

apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: ExternalName
  spec:
    externalName: example.com

```

**ExternalName:**  
Maps the Service to the contents of the ExternalName field (e.g. example.com), by returning a CNAME record with its value.

**Example Usecases:**  
When you expose and stick your traffic directly to your subdomain

CLARUSWAY®  
WAY TO REINVENT YOURSELF

48